Pilot training runs: shielded agent, tuning seed 100, base level (noise 0.1)

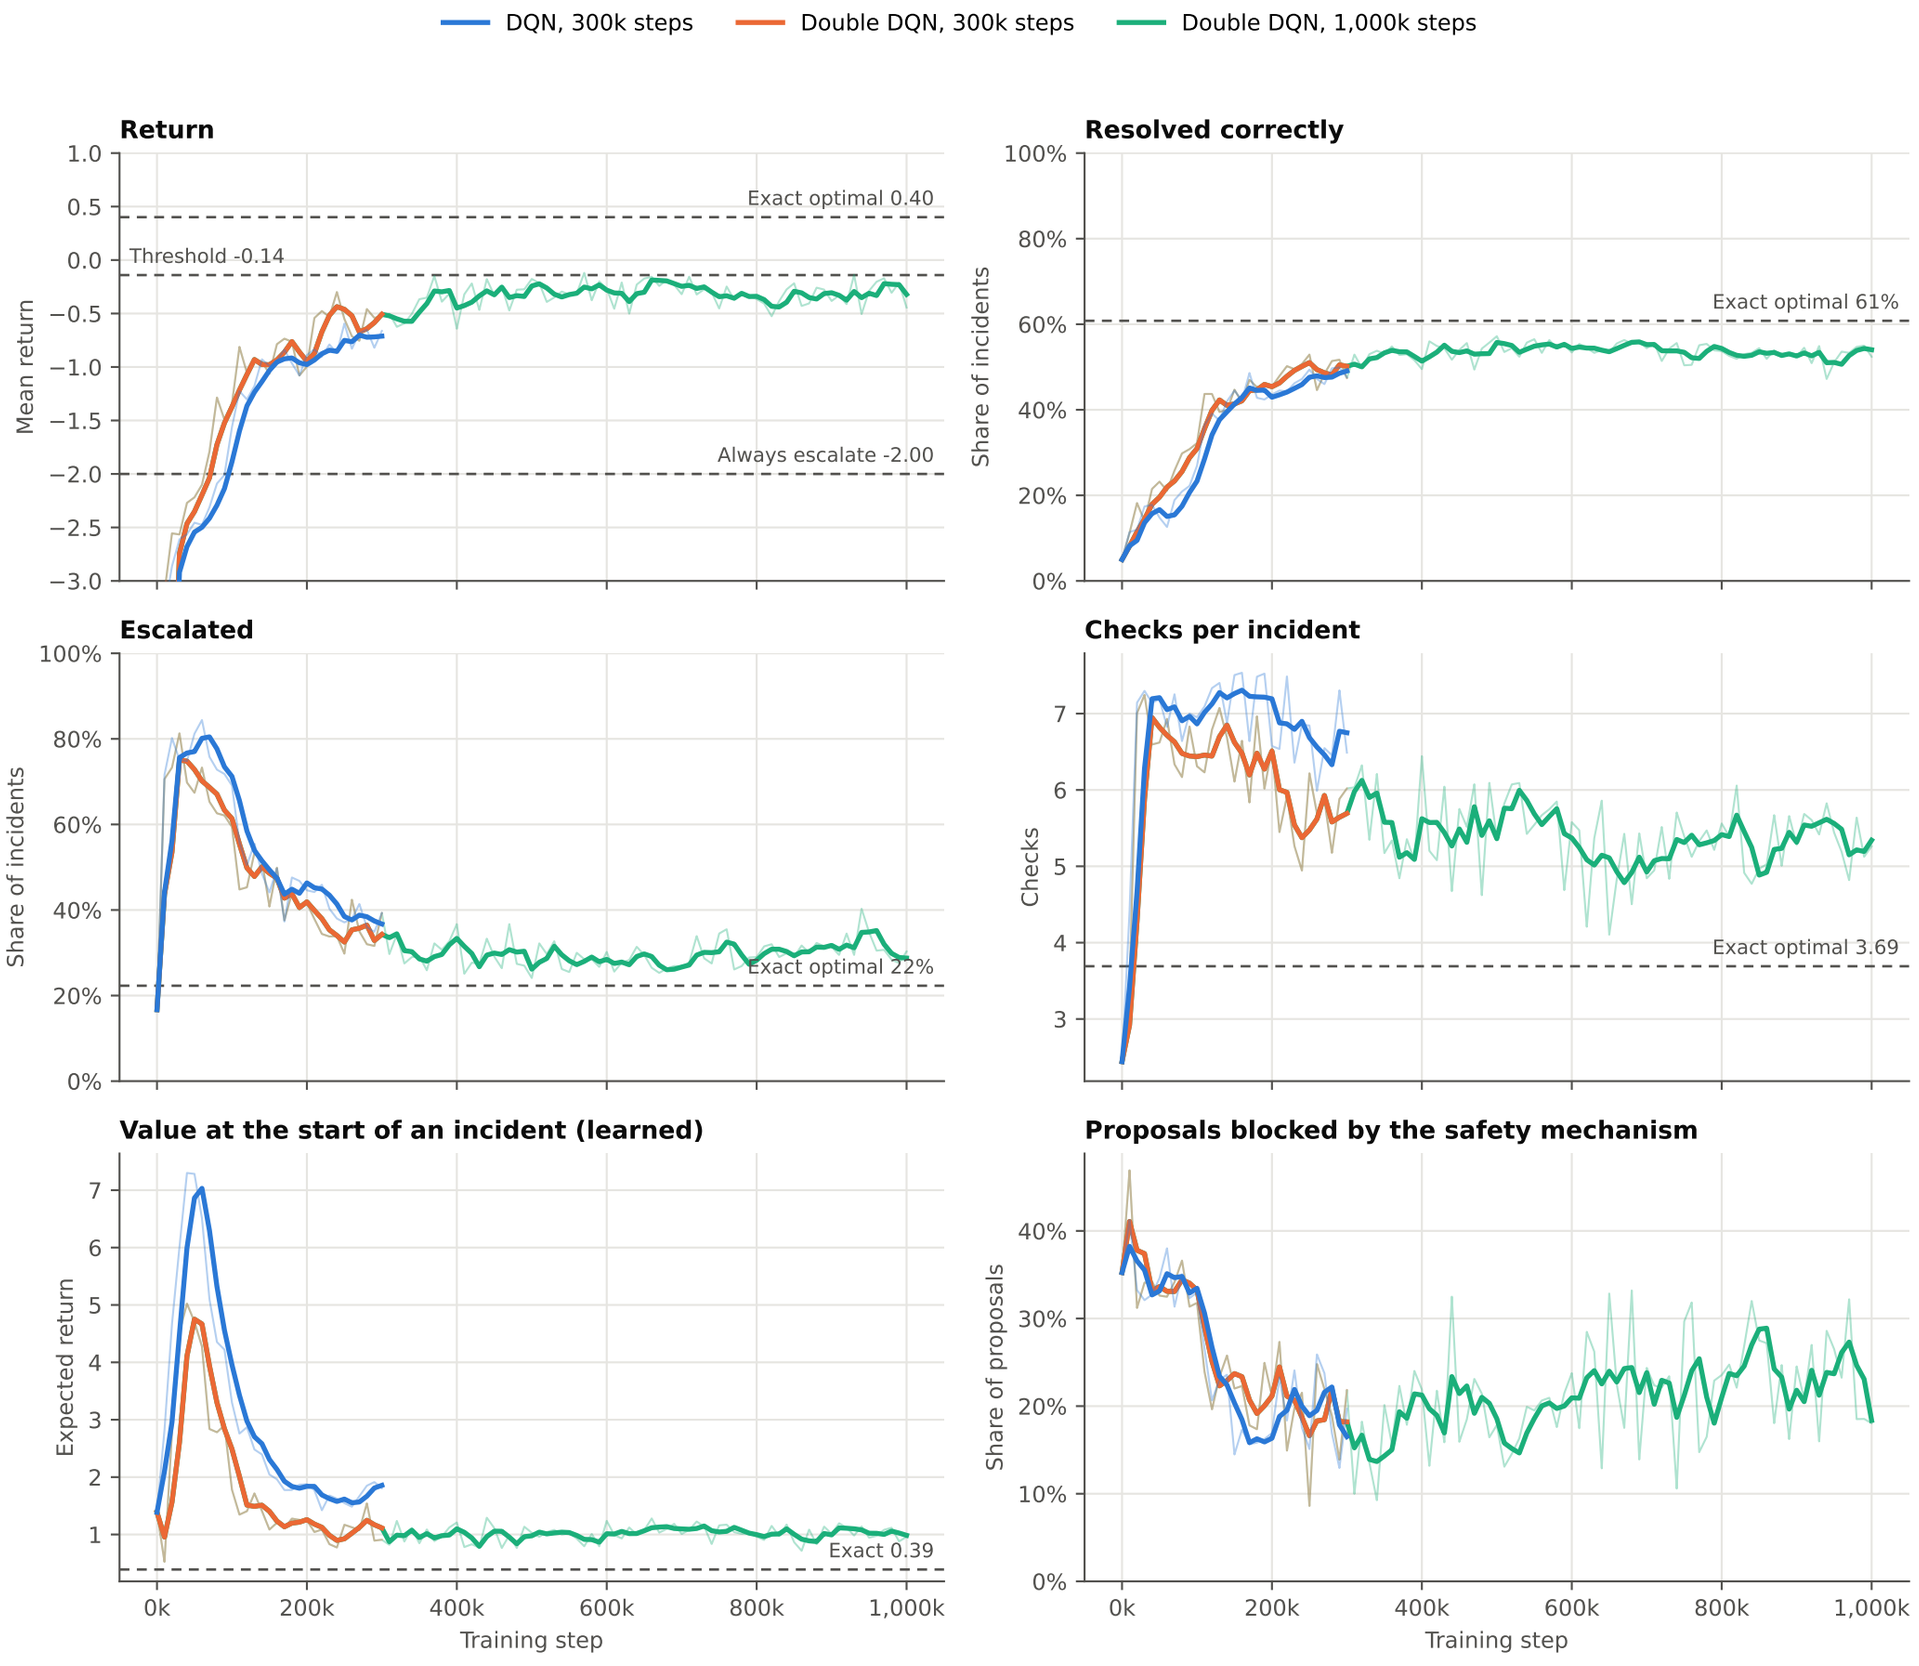

Each point: 1,000 monitoring incidents, no exploration. Thin lines: each evaluation; thick: rolling mean of three. Return panel starts at -3; the first evaluations (about -8) are off scale. Pilot data (one developer's ratings), not Directorate data. Double DQN at 300k and at 1M steps use the same seed and settings and are identical for their first 300k steps, so the 300k line lies on top of the 1M line there.