Training the three agents with fix-and-verify (Experiment 2 runs, 3 seeds each)

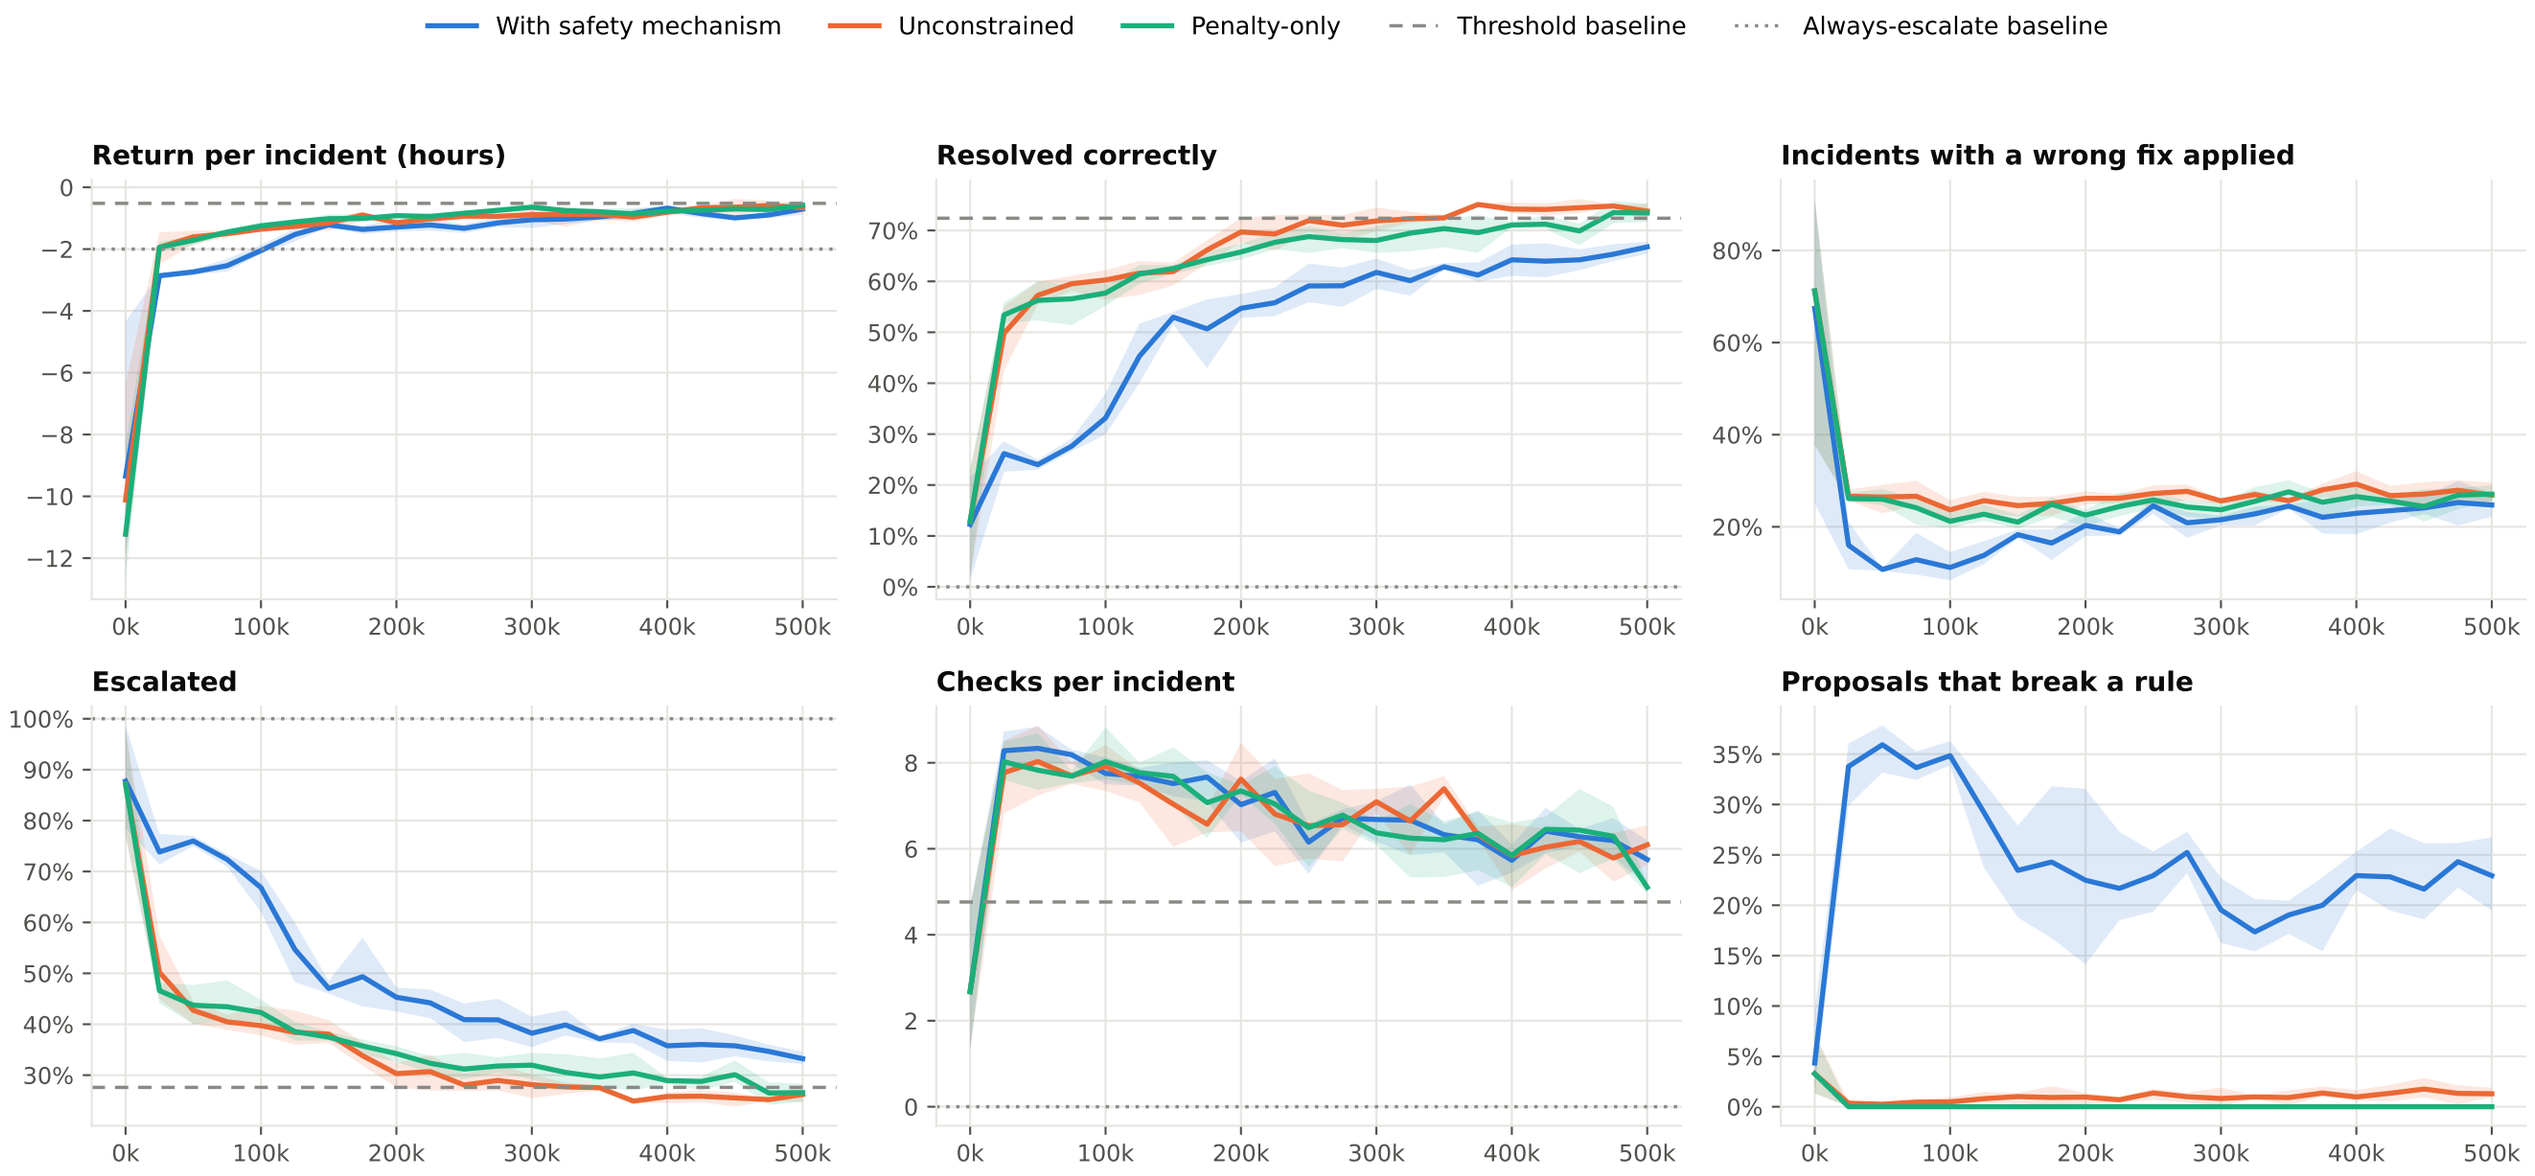

Training monitor: 1,000 fixed monitoring incidents every 25k steps, greedy policy (not the evaluation incidents). Line = mean over seeds; band = lowest to highest seed. Model v2 (pilot developer, Part 2) with fix-with-verification settings; the verification values are researcher-provisional. Pilot data, not Directorate data. Baselines: the same monitoring incidents, with the safety mechanism.